

Compute unit in  
`objects_detector_p->compute();`

In function

`void ObjectsDetectionApplication::one_step(const cv::Mat& m_detecting_image)`

We assume the `objects_detector_p` has been init

## GpuIntegralChannelsDetector.cpp

```
void GpuIntegralChannelsDetector::compute()
{
    detections.clear();
    num_gpu_detections = 0; // no need to clean the buffer

    // some debugging variables
    static bool first_call = true;

    assert(integral_channels_computer_p);
    //assert(gpu_detection_variant_cascade_per_scale.getBuffer() != NULL);

    // for each range search
    for(size_t search_range_index = 0; search_range_index < search_ranges_data.size(); search_range_index +=1)
    {
        compute_detections_at_specific_scale_v1(search_range_index, first_call);
    } // end of "for each search range"

    collect_the_gpu_detections();

    log.info() << "number of raw (before non maximal suppression) detections on this frame == "
        << detections.size() << std::endl;

    // windows size adjustment should be done before non-maximal suppression

    (*model_window_to_object_window_converter_p)(detections);

    compute_non_maximal_suppression();

    first_call = false;

    return;
}
```

Search range has been initialized,  
so what's the **search range** mean?  
\\src\\objects\_detection\\DetectorSearchRange.hpp

Page 6 `compute_detections_at_specific_scale_v1(search_range_index, first_call);`

Page 10 `(*model_window_to_object_window_converter_p)(detections);`

`compute_non_maximal_suppression();`

## IntegralChannelsDetector.cpp

```
void IntegralChannelsDetector::compute()
{
    detections.clear();

    // some debugging variables
    static bool first_call = true;

    assert(integral_channels_computer_p);
    assert(search_ranges_data.size() == detection_cascade_per_scale.size());

    // for each range search
    for(size_t search_range_index=0; search_range_index < search_ranges_data.size(); search_range_index +=1)
    {
        compute_detections_at_specific_scale(search_range_index, first_call);
    } // end of "for each search range"

    process_raw_detections();

    first_call = false;

    return;
}

void IntegralChannelsDetector::process_raw_detections()
{
    const size_t num_raw_detections = detections.size();

    // windows size adjustment should be done before non-maximal suppression

    (*model_window_to_object_window_converter_p)(detections);

    log_info() << "number of detections (before non maximal suppression) on this frame == "
        << num_raw_detections << " (raw) / " << detections.size() << " (after filtering)" << std::endl;

    compute_non_maximal_suppression();

    return;
}
```

First, we focus on the key step, compute.

In compute both cpu/gpu version call a function like this.

```
compute_detections_at_specific_scale(search_range_index, first_call);
```

so what's the **search range** mean?

In next page, I crop the output of the example program using trained model **headhunter** with the config file eccv2014 face detection pascal.config.ini



22 unique Integral Channel components  
With 2 scales, 11 components(11  
semantic categories) each scale.

In each **scale** to be searched, choose a  
**scales** near it and generate a search range  
for each **component**  
One search range bind to one certain detector  
component and one scale of the input image

```
x_stride = 0.001  
y_stride = 0.001  
  
non_maximal_suppression_method = greedy  
minimal_overlap_threshold=0.3  
  
min_scale = 0.325  
max_scale = 6  
  
num_scales = 30
```

The input image will be resize to 30  
scales

- In this example we will have  $30 \times 11 = 330$  search range.

[illegible]

11 components  
each scale

## Search range

Consider one search range in last slide.

## GpuIntegralChannelsDetector.cpp

```
void GpuIntegralChannelsDetector::compute_detections_at_specific_scale_v1(  
    const size_t search_range_index,  
    const bool first_call)  
{
```

```
    doppia::objects_detection::gpu_integral_channels_t &integral_channels =  
        resize_input_and_compute_integral_channels(search_range_index, first_call);
```

```
    const ScaleData &scale_data = extra_data_per_scale[search_range_index];
```

```
    // compute the scores --  
    invoke_v1_integral_channels_detector visitor(  
        integral_channels,  
        search_range_index,  
        scale_data,  
        score_threshold,  
        gpu_detections, num_gpu_detections);
```


```
    // compute the detections, and keep the results on the gpu memory  
    boost::apply_visitor(visitor, gpu_detection_variant_cascade_per_scale);
```

```
    // ( the detections will be collected after iterating over all the scales )
```

```
#if defined(BOOTSTRAPPING_LIB)  
    throw std::runtime_error("GpuIntegralChannelsDetector::compute_detections_at_specific_scale_v1 "  
        "should not be used inside bootstrapping_lib, use v0 instead");  
#endif  
    return;  
}
```

- Resize the input image to the scale.
- Computer the integral feature of the image at that scale  
(to speed up we can compute only once in each scale but not done here)

- Search faces in the search range above.



10 channels =  
6 hog features  
1 gradient  
3 LUV (color)

This is the 10 channels  
and integral channels are the integral of these channels.

## \\src\\objects\_detection\\gpu\\integral\_channels\_detector.cu

```
671 /// this method directly adds elements into the gpu_detections vector
672 template<typename GpuDetectionCascadePerScaleType>
673 void integral_channels_detector_impl(gpu_integral_channels_t &integral_channels,
674                                     const size_t search_range_index,
675                                     const doppia::ScaleData &scale_data,
676                                     GpuDetectionCascadePerScaleType &detection_cascade_per_scale,
677                                     const float score_threshold,
678                                     gpu_detections_t& gpu_detections,
679                                     size_t &num_detections)
680 {
    :
    :
    :
724 const int
725     width = search_range.max_x - search_range.min_x,
726     height = search_range.max_y - search_range.min_y;
727
728 if((width <= 0) or (height <= 0))
729 { // nothing to be done
730     // num_detections is left unchanged
731     return;
732 }
733
734 dim3 grid_dimensions(div_up(width, block_dimensions.x),
735                      div_up(height, block_dimensions.y));
736
737 // prepare variables for kernel call --
738 bind_integral_channels_texture(integral_channels);
739 move_num_detections_from_cpu_to_gpu(num_detections);
740
741 integral_channels_detector_kernel
742     <CascadeStageType>
743     <<<grid_dimensions, block_dimensions>>>
744         (scale_datum,
745          integral_channels,
746          search_range_index,
747          detection_cascade_per_scale,
748          score_threshold,
749          gpu_detections);
750
```

- Search faces in the search range
- This function compute the each possible windows parallel.
- The detector (the cascade) Compute the score of the windows and decide if it is a face with the score\_threshold defined in config file.

```
20 score_threshold = 0.05 # nice for visualization
21 #score_threshold = 0 # default threshold
```

# src\objects\_detection\gpu\integral\_channels\_detector.cu

```

3/// this kernel is called for each position where we which to detect objects
/// we assume that the border effects where already checked when computing the DetectorSearchRange
/// thus we do not do any checks here.
/// This kernel is a mirror of the CPU method compute_cascade_stage_on_row(...) inside IntegralChannelsDetector.cpp
/// @see IntegralChannelsDetector
template <typename DetectionCascadeStageType>
__global__
void integral_channels_detector_kernel(
    const gpu_scale_datum_t scale_datum,
    const gpu_integral_channels_t::KernelConstData integral_channels,
    const size_t scale_index,
    const typename Cuda::DeviceMemory<DetectionCascadeStageType, 2>::KernelConstData detection_cascade_per_scale,
    const float score_threshold,
    gpu_detections_t::KernelData gpu_detections)
{
    :
    :
    :

```

```

// retrieve current score value
float detection_score = 0;

```

```

const size_t
    cascade_length = detection_cascade_per_scale.size[0],
    scale_offset = scale_index * detection_cascade_per_scale.stride[0];

```

```

for(size_t stage_index = 0; stage_index < cascade_length; stage_index += 1)
{
    const size_t index = scale_offset + stage_index;

    // we copy the cascade stage from global memory to thread memory
    // (when using a reference code runs at ~4.35 Hz, with copy it runs at ~4.55 Hz)
    const DetectionCascadeStageType stage = detection_cascade_per_scale.data[index];

    update_detection_score(x, y, stage, integral_channels, detection_score);

    if(detection_score < stage.cascade_threshold)
    {
        // this is not an object of the class we are looking for
        // do an early stop of this pixel
        detection_score = -1E5; // since re-ordered classifiers may have a "very high threshold"
        break;
    }
} // end of "for each stage"

```

```

// >= to be consistent with Markus's code
if(detection_score >= score_threshold)
{
    // we got a detection
    add_detection(gpu_detections, x, y, scale_index, detection_score);
}

return;

```

- The final score is the sum of all stages in a cascade

```

stages {
    feature_type: Level2DecisionTree
    weight: 0.829010546207
    cascade_threshold: -inf
}

nodes {
    id: 0
    parent_id: 0
    parent_value: true
    decision_stump {
        feature {
            channel_index: 3
            box {
                min_corner {
                    x: 18
                    y: 25
                }
                max_corner {
                    x: 22
                    y: 29
                }
            }
        }
        feature_threshold: 49.0
        larger_than_threshold: false
    }
}

```

Never break, sum up all the stages in this case

```

nodes {
    id: 1
    parent_id: 0
    parent_value: true
    decision_stump {
        feature {
            channel_index: 7
            box {
                min_corner {
                    x: 0
                    y: 29
                }
                max_corner {
                    x: 1
                    y: 30
                }
            }
        }
        feature_threshold: 1.0
        larger_than_threshold: false
    }
}

```

```

nodes {
    id: 2
    parent_id: 0
    parent_value: false
    decision_stump {
        feature {
            channel_index: 3
            box {
                min_corner {
                    x: 5
                    y: 22
                }
                max_corner {
                    x: 13
                    y: 30
                }
            }
        }
        feature_threshold: 91.0
        larger_than_threshold: false
    }
}

```



# src\objects\_detection\gpu\integral\_channels\_detector.cu

In root stump, score compared with threshold return true or false.

In leaf stump, Score compared with threshold return +weight or -weight

```

256 const float level1_feature_value =
257     get_feature_value(feature_t, use_2d_texture)(
258         weak_classifier.level1_node.feature, x, y, integral_channels);
284 if (evaluate_decision_stump(weak_classifier.level1_node, level1_feature_value))
285 {
286     const float level2_true_feature_value =
287         get_feature_value(feature_t, use_2d_texture)(
288             weak_classifier.level2_true_node.feature, x, y, integral_channels);
289     current_score += evaluate_decision_stump(weak_classifier.level2_true_node, level2_true_feature_value);
291 }
292 else
293 {
294     const float level2_false_feature_value =
295         get_feature_value(feature_t, use_2d_texture)(
296             weak_classifier.level2_false_node.feature, x, y, integral_channels);
297     current_score += evaluate_decision_stump(weak_classifier.level2_false_node, level2_false_feature_value);
298 }
299

```

```

stages {
    feature_type: Level2DecisionTree
    weight: 0.829010546207
    cascade_threshold: -inf

```


```

nodes {
    id: 0
    parent_id: 0
    parent_value: true
    decision_stump {
        feature {
            channel_index: 3
            box {
                min_corner {
                    x: 18
                    y: 25
                }
                max_corner {
                    x: 22
                    y: 29
                }
            }
        }
        feature_threshold: 49.0
        larger_than_threshold: false
    }
}

```

## feature

Dollár, Piotr, et al. "Integral Channel Features." *BMVC*. Vol. 2. No. 3. 2009.  
[http://vision.ucsd.edu/sites/default/files/dollarBMVC09ChnFtrs\\_0.pdf](http://vision.ucsd.edu/sites/default/files/dollarBMVC09ChnFtrs_0.pdf)



10 channels =  
 6 hog features  
 1 gradient  
 3 LUV (color)


```

feature {
    channel_index: 4
    box {
        min_corner {
            x: 17
            y: 26
        }
        max_corner {
            x: 21
            y: 30
        }
    }
}

```

Score = sum of the value  
 in channel n  
 within the box.

1. For one input picture,  
 calculate integral image on each  
 channels. Which can be reused  
 when calculate the feature.
2. Calculate score in one feature  
 box.
3. Compare to the threshold,  
 output true or false.



```

nodes {
    id: 1
    parent_id: 0
    parent_value: true
    decision_stump {
        feature {
            channel_index: 7
            box {
                min_corner {
                    x: 0
                    y: 29
                }
                max_corner {
                    x: 1
                    y: 30
                }
            }
        }
        feature_threshold: 1.0
        larger_than_threshold: false
    }
}

```

```

nodes {
    id: 2
    parent_id: 0
    parent_value: false
    decision_stump {
        feature {
            channel_index: 3
            box {
                min_corner {
                    x: 5
                    y: 22
                }
                max_corner {
                    x: 13
                    y: 30
                }
            }
        }
        feature_threshold: 91.0
        larger_than_threshold: false
    }
}

```

# objects\_detection\non\_maximal\_suppression\GreedyNonMaximalSuppression.cpp

```
275 void GreedyNonMaximalSuppression::compute()
276 {
277     candidate_detections.sort(has_higher_score);
278     maximal_detections.clear();
279     maximal_detections.reserve(64); // we do not expect more than 64 pedestrians per scene
280
281     candidate_detections_t::iterator detections_it = candidate_detections.begin();
282     for(; detections_it != candidate_detections.end(); ++detections_it)
283     {
284         const detection_t &detection = *detections_it;
285
286         // this detection passed the test
287         maximal_detections.push_back(detection);
288
289         candidate_detections_t::iterator lower_score_detection_it = detections_it;
290         ++lower_score_detection_it; // = detections_it + 1
291         for(; lower_score_detection_it != candidate_detections.end(); )
292         {
293             const float overlap = compute_overlap_inlined(detection, *lower_score_detection_it, overlap_method);
294
295             if(overlap > minimal_overlap_threshold)
296             {
297                 // this detection seems to overlap too much, we should remove it
298                 lower_score_detection_it = candidate_detections.erase(lower_score_detection_it);
299             }
300             else
301             {
302                 // we keep this detection in the candidates list
303                 ++lower_score_detection_it;
304             }
305
306         } // end of "for each lower score detection"
307
308     } // end of "for each candidate detection"
309
310     return;
311 }
312 }
```

- Sort the candidate by score
- remove the overlap candidate with small score

# Summary for the detection

